



input layer  
(states ( $\mathbf{s}$ ))

hidden layer  
(policy ( $\pi_\theta$ ))

output layer  
(actions ( $\mathbf{a}$ ))



policy parameterized  
with a neural network